

ASSIGNMENT 6: LOGISTIC REGRESSION



Markus Holzleitner,
Johannes Kofler,
Angela Bitto-Nemling
Institute for Machine Learning

Copyright statement:

This material, no matter whether in printed or electronic form, may be used for personal and non-commercial educational use only. Any reproduction of this material, no matter whether as a whole or in parts, no matter whether in printed or in electronic form, requires explicit prior acceptance of the authors.

Logistic Regression

- Given a bunch of data $\mathbf{Z} = \{(\mathbf{x}_1, \mathbf{y}_1), (\mathbf{x}_2, \mathbf{y}_2), \dots, (\mathbf{x}_l, \mathbf{y}_l)\}$

- ☐ \mathbf{x}_i ... vectors of features

- ☐ \mathbf{y}_i ... labels

- **Classification:** $y_i \in \{1, \dots, M\}$

- **First (bad) idea:** fit a linear regression line


$$g_{LR}(\mathbf{x}_i) = \mathbf{W}^T \mathbf{x}_i$$

- **Goal:**

$$y_i = \begin{cases} 0 & g_{LR}(\mathbf{x}_i) < 0.5 \\ 1 & g_{LR}(\mathbf{x}_i) \geq 0.5 \end{cases}$$

- **Problem:** the resulting $y_i(\mathbf{x}_i)$ would be a step function, i.e. not continuous

Problem with Linear Regression





Logistic Regression

- Instead of predicting directly the class we would like to predict the class probability, which is a regression problem
- Therefore we apply the logistic function:



$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

- z is a linear function of features: $z = \mathbf{W}^T \mathbf{x}$
- Model:

$$g(\mathbf{x}) = \sigma(\mathbf{W}^T \mathbf{x})$$




Logistic Function



Also known as **sigmoid function**.

Also known as **Fermi function** in physics.

Logistic Regression





Objective

- Likelihood function for a Bernoulli distribution:



$$\mathcal{L}(\{\mathbf{x}, y\}; \mathbf{W}) = \prod_{i=1}^n g(\mathbf{x}_i; \mathbf{W})^{y_i} \cdot (1 - g(\mathbf{x}_i; \mathbf{W}))^{1-y_i}$$

- Taking the negative logarithm, we obtain:

$$\begin{aligned} L &= -\log \mathcal{L}(\{\mathbf{x}, y\}; \mathbf{W}) \\ &= -\sum_i \left(y_i \log g(\mathbf{x}_i; \mathbf{W}) + (1 - y_i) \log(1 - g(\mathbf{x}_i; \mathbf{W})) \right) \end{aligned}$$

- L is known as the **Cross Entropy Loss**
- The task $\min_{\mathbf{W}} L$ does in general not have a closed-form solution
- Fortunately L is a **convex function** for Logistic Regression, therefore, every local minimum is a global minimum.

Convex vs. non-convex





Gradients for Logistic Regression

- **Using:** $\frac{\partial \sigma(z)}{\partial z} = \sigma(z) \cdot (1 - \sigma(z))$
- **Using** σ_i instead of $\sigma(\mathbf{W}^T \mathbf{x}_i)$

$$\begin{aligned} L &= - \sum_i \left(y_i \log \sigma(\mathbf{W}^T \mathbf{x}_i) + (1 - y_i) \log(1 - \sigma(\mathbf{W}^T \mathbf{x}_i)) \right) \\ \frac{\partial L}{\partial \mathbf{W}} &= - \sum_i \left(y_i \frac{1}{\sigma_i} \cdot \sigma_i \cdot (1 - \sigma_i) \cdot \mathbf{x}_i - \frac{1 - y_i}{1 - \sigma_i} \cdot \sigma_i \cdot (1 - \sigma_i) \cdot \mathbf{x}_i \right) \\ &= - \sum_i \left(y_i(1 - \sigma_i) \cdot \mathbf{x}_i - (1 - y_i) \cdot \sigma_i \cdot \mathbf{x}_i \right) \\ &= - \sum_i (y_i - y_i \sigma_i - \sigma_i + \sigma_i y_i) \mathbf{x}_i \\ &= \sum_i (\sigma_i - y_i) \mathbf{x}_i \end{aligned}$$



Logistic Regression Problem

■ Task:

$$\min_{\mathbf{W}} L = \min_{\mathbf{W}} \left[- \sum_i (y_i \log g(\mathbf{x}^i; \mathbf{W}) + (1 - y_i) \log(1 - g(\mathbf{x}^i; \mathbf{W}))) \right]$$

$$\square g(\mathbf{x}; \mathbf{W}) = \sigma(\mathbf{W}^T \mathbf{x}) = \frac{1}{1 + e^{-\mathbf{W}^T \mathbf{x}}}$$

■ Note: no closed-form solution!

You have to use methods like [Gradient Descent](#), Newton, BFGS, Conjugate Gradient, ...



Gradient Descent

- **Given:** a function $f(x)$
- **Task:** find x that maximizes (or minimizes) $f(x)$
- **Idea:** start at some value x_0 , and take a small step η in the direction in which the function decreases strongest.

Find derivative $f' = \frac{\partial f}{\partial x}$.

- $-f'(x_0)$ is the direction of steepest descent at x_0 .

- **Solution:**

- Iteratively calculate $x_{i+1} = x_i - \eta \cdot f'(x_i)$.
- Each x_{i+1} should be a better solution than x_i .
- Eventually you'll reach a (**local**) minimum.

Gradient descent

Gradient descent for non-convex problems

Gradient Descent in Logistic Regression



The minimization of the loss function $L(., \mathbf{W})$ can be done by Gradient Descent:

$$\mathbf{W}_{n+1} = \mathbf{W}_n - \eta \frac{\partial L}{\partial \mathbf{W}} ,$$

where η is the learning rate,

and \mathbf{W}_0 is some initial guess for \mathbf{W} .



Gradient Checking

- Method for checking if the symbolic computation/implementation of the gradient was correct.
- Logistic Regression gradient is easy, but once we get to neural networks, you'll be glad to know this trick.
- **Idea:** compare your gradient with a numerical approximation of the gradient.



Gradient Checking

- Central difference quotient:

$$\frac{\partial L}{\partial W_{ij}} \approx \frac{L(., \mathbf{W} + \epsilon \mathbf{e}_{ij}) - L(., \mathbf{W} - \epsilon \mathbf{e}_{ij})}{2 \epsilon}$$

- Central difference quotient for logistic regression (\mathbf{W} is a vector):

$$\frac{\partial L}{\partial W_i} \approx \frac{L(., \mathbf{W} + \epsilon \mathbf{e}_i) - L(., \mathbf{W} - \epsilon \mathbf{e}_i)}{2 \epsilon}$$

with $\mathbf{e}_i = (0 \ 0 \ \dots \ 1 \ \dots \ 0)^T$.

- Good choice is $\epsilon = 10^{-4}$.